



Swinburne University of Technology Hawthorn Campus
Department of Computing Technologies

COS20028 Big Data Architecture and Application
Assignment 2 - Semester 2, 2023

Name: [Nguyen Nam Tung](#)

Student ID: [103181157](#)

Submission Date: [26/10/2023](#)

Assignment Task

1. Find a way with tools taught in this unit to check whether the language_code attribute in the source data is unique. **Note that you cannot eyeball the result. The result should be found by combining valid tools in a sequence.**
 - 1.1. List the sequence of tools you used for finding the answer:
[Ans: Hive can be utilized to find the answers.](#)
 - 1.2. Give a short explanation of which tool did what.
[Ans: Hive counts the number of all the language_code attribute.](#)
[If the count is greater than 1 then it is not unique.](#)
 - 1.3. List the code/command/statement and the outcome screenshot of the step in the sequence:
(Note that only meaningful screenshot is required. For example, the screenshot of the MapReduce execution result is necessary, but the screenshot of the outcome from the "cd .." command is not meaningful)

[Phase 1 – Querying to check that the language_code attribute is unique]**Query:**

```
SELECT lng_code, COUNT(lng_code) FROM austlang GROUP BY lng_code
HAVING COUNT(lng_code) > 1;
```

Outcome:**No Data Available for this query so that language code is unique**

```
hive> SELECT lng_code, COUNT(lng_code) FROM austlang GROUP BY lng_code HAVING COUNT(lng_code) > 1;
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_202310261911_0007, Tracking URL = http://0.0.0.0:50030/jobdetails.jsp?jobid=job_202310261911_0007
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_202310261911_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-10-26 20:25:21,891 Stage-1 map = 0%, reduce = 0%
2023-10-26 20:25:23,898 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.03 sec
2023-10-26 20:25:24,905 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.03 sec
2023-10-26 20:25:25,911 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.03 sec
2023-10-26 20:25:26,918 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.57 sec
2023-10-26 20:25:27,926 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.57 sec
2023-10-26 20:25:28,933 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.57 sec
MapReduce Total cumulative CPU time: 2 seconds 570 msec
Ended Job = job_202310261911_0007
MapReduce Jobs Launched:
Job 0: Map: 1 Reduce: 1 Cumulative CPU: 2.57 sec HDFS Read: 320819 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 570 msec
JK
Time taken: 11.008 seconds
hive>
```

training@localhost:~

2. Find a way with tools taught in this unit to list the unique values for all entities (tables with the solid boundary) except lng_id.
 - 2.1. What is the best tool of choice to perform this task?
Ans: MapReduce is utilized to perform this task.
 - 2.2. Assume you choose MapReduce to be the tool. Should it be the Map-only design or the complete MapReduce with mapper and reducer involved in the process?
Ans: No, it should not be the Map – only design. It should have both Mapper and Reducer as we need a MapReduce design to aggregate the data to find the unique values.
 - 2.3. Assume you choose MapReduce to be the tool. Does the count in the output matter for preparing the data for creating the entities?
Ans: No because we only want to list the unique values for all the entities. But for example, in question 2.5, the count in the output might be mattered to find the counts of a particular state.
 - 2.4. List the code/command/statement and the outcome screenshot of preparing the data for lng_st.

Mapper Code:

```
package stubs;

import java.io.IOException;

public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {

    private Text Key = new Text();
    private IntWritable Value = new IntWritable();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        String line = value.toString();
        String[] data = line.split("\\t");
        String[] data_lng_st = data[7].split(",");
        for (String data_split : data_lng_st){
            if (data_split.length() > 0){
                Key.set(data_split);
                Value.set(1);
                context.write(Key, Value);
            }
        }
    }
}
```

Reducer Code:

```
package stubs;

import java.io.IOException;

public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    private IntWritable Value = new IntWritable();

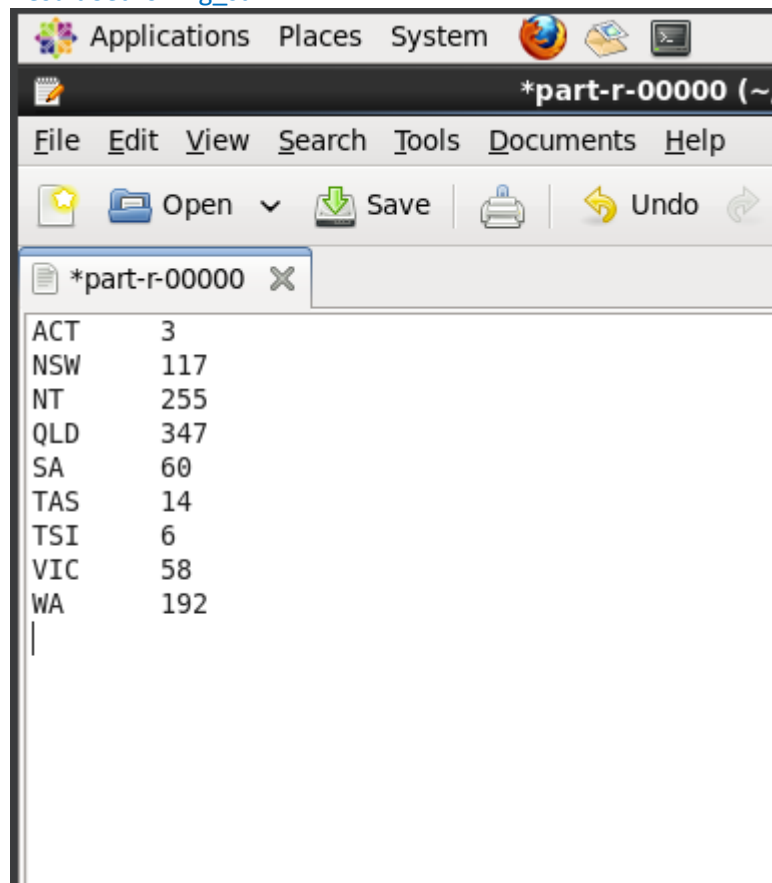
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int count = 0;

        for (IntWritable value : values){
            count += 1;
        }

        Value.set(count);
        context.write(key, Value);
    }
}
```

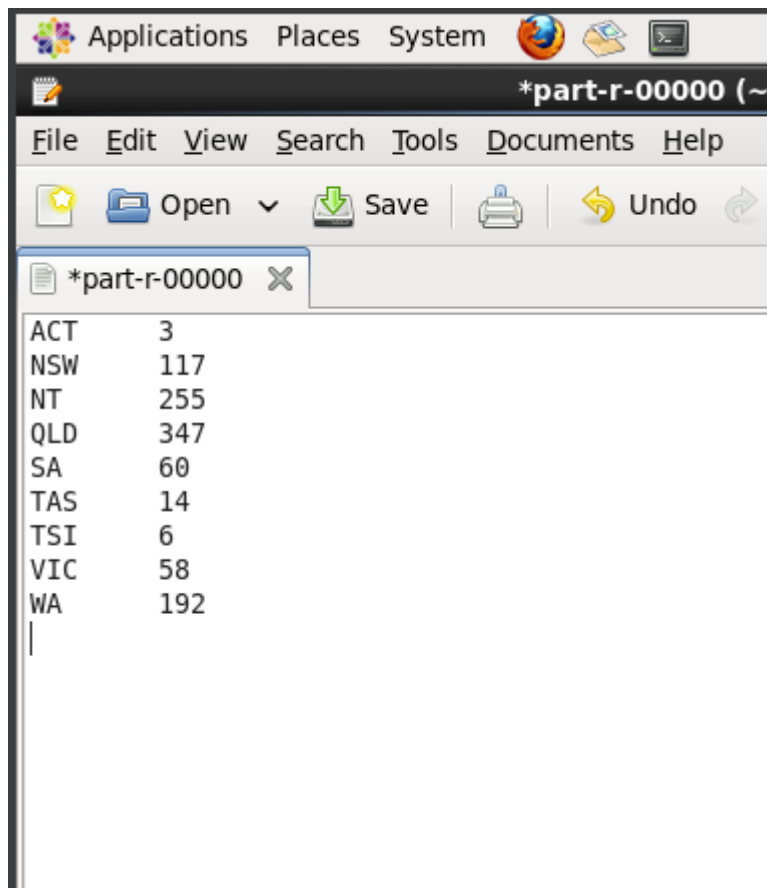
Result Set for lng_st:



ACT	3
NSW	117
NT	255
QLD	347
SA	60
TAS	14
TSI	6
VIC	58
WA	192

2.5. How many counts does QLD have?

[We can check the output file for the desired answer.](#)



The screenshot shows a text editor window titled '*part-r-00000 (~/' with a menu bar (File, Edit, View, Search, Tools, Documents, Help) and a toolbar (Open, Save, Undo). The editor contains a table with two columns: state abbreviations and their corresponding counts.

ACT	3
NSW	117
NT	255
QLD	347
SA	60
TAS	14
TSI	6
VIC	58
WA	192

Ans: 347

3. Find a way with tools taught in this unit to prepare the data for all entities (tables with the solid boundary) for MySQL and Hive.

- 3.1. Assume your tool of choice is MapReduce for this task. Will this be a Map-only job or a complete MapReduce job for preparing the desired data?

Ans: In my opinion, it should be a full MapReduce job. This is because, we need to aggregate the data to make sure that a unique key is generated for each record. However, on the other hand, a map-only job may be enough to get the records from particular fields from the original dataset.

- 3.2. List the code/command/statement and the outcome screenshot of preparing the data for lng_id.

Mapper Code:

```
package stubs;

import java.io.IOException;

public class WordMapper extends Mapper<LongWritable, Text, Text, Text> {

    private Text Key = new Text();
    private Text Key2 = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString();
        String[] data = line.split("\\t");
        Key.set(data[0] + '\t' + data[5] + '\t' + data[6] + '\t' + data[8]);
        Key2.set("");
        context.write(Key, Key2);
    }
}
```

Reducer Code:

```
package stubs;

import java.io.IOException;

public class SumReducer extends Reducer<Text, IntWritable, Text, Text> {

    private Text Value = new Text();

    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        Value.set("");
        context.write(key, Value);
    }
}
```

Result Set for lng_id

S/A1	-32.39094519	118.7550827	https://collection.aiatsis.gov.au/austlang/language/a1
A10			https://collection.aiatsis.gov.au/austlang/language/a10
A100			https://collection.aiatsis.gov.au/austlang/language/a100
A101	-28.1355411	114.7644712	https://collection.aiatsis.gov.au/austlang/language/a101
A102	-29.46595664	127.7662611	https://collection.aiatsis.gov.au/austlang/language/a102
A103	-29.86593028	122.698712	https://collection.aiatsis.gov.au/austlang/language/a103
A104			https://collection.aiatsis.gov.au/austlang/language/a104
A105			https://collection.aiatsis.gov.au/austlang/language/a105
A106			https://collection.aiatsis.gov.au/austlang/language/a106
A107			https://collection.aiatsis.gov.au/austlang/language/a107
A108			https://collection.aiatsis.gov.au/austlang/language/a108
A109			https://collection.aiatsis.gov.au/austlang/language/a109
A11	-29.231154	122.7229745	https://collection.aiatsis.gov.au/austlang/language/a11
A110			https://collection.aiatsis.gov.au/austlang/language/a110
A111			https://collection.aiatsis.gov.au/austlang/language/a111
A112			https://collection.aiatsis.gov.au/austlang/language/a112
A113			https://collection.aiatsis.gov.au/austlang/language/a113
A114			https://collection.aiatsis.gov.au/austlang/language/a114
A115			https://collection.aiatsis.gov.au/austlang/language/a115
A116			https://collection.aiatsis.gov.au/austlang/language/a116
A117			https://collection.aiatsis.gov.au/austlang/language/a117
A118			https://collection.aiatsis.gov.au/austlang/language/a118
A119			https://collection.aiatsis.gov.au/austlang/language/a119
A12	-29.77825242	121.9180727	https://collection.aiatsis.gov.au/austlang/language/a12
A120			https://collection.aiatsis.gov.au/austlang/language/a120
A121			https://collection.aiatsis.gov.au/austlang/language/a121
A122			https://collection.aiatsis.gov.au/austlang/language/a122
A123			https://collection.aiatsis.gov.au/austlang/language/a123
A124			https://collection.aiatsis.gov.au/austlang/language/a124
A13	-28.82675757	116.721527	https://collection.aiatsis.gov.au/austlang/language/a13
A14	-28.59330827	117.1853249	https://collection.aiatsis.gov.au/austlang/language/a14
A16	-28.1820317	120.3453645	https://collection.aiatsis.gov.au/austlang/language/a16
A17	-28.68370166	124.8034463	https://collection.aiatsis.gov.au/austlang/language/a17

4. Find a way with tools taught in this unit to prepare the data for all weak entities (tables with the dashed boundary) for MySQL and Hive.

- 4.1. Which tool would be the best choice for handling this task?

Ans: MapReduce can be utilized to handle this task.

- 4.2. List the code/statement of preparing the data for all weak entities.

Mapper Code:

```
package stubs;

import java.io.IOException;

public class WordMapper extends Mapper<LongWritable, Text, Text, Text> {

    private Text Key = new Text();
    private Text Key2 = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString();
        String[] data = line.split("\\t");
        String lng_name = data[1];
        String lng_code = data[0];
        String[] lng_name_part = lng_name.split("/");
        if (lng_name_part.length > 0){
            for (String name : lng_name_part){
                context.write(new Text(lng_code), new Text(name));
            }
        }
    }
}
```

4.3. Post the first-page screenshot of the rel_code_name processed result.

Ans:



5. Follow the given ERDs, create tables and import data into the created tables. All tables should be put in the database entitled “indigenous”.

5.1. List all statements of this task for MySQL:

```
Database: indigenous
mysql> CREATE TABLE lng_name (lng_name VARCHAR(60) NOT NULL, PRIMARY KEY (lng_name));
Query OK, 0 rows affected (0.02 sec)

mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/lng_name/lng_name.txt' INTO TABLE lng_name FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 1243 rows affected (0.00 sec)
Records: 1243 Deleted: 0 Skipped: 0 Warnings: 0

mysql> CREATE TABLE lng_synonym (lng_synonym VARCHAR(1000) , PRIMARY KEY (lng_synonym));
Query OK, 0 rows affected (0.00 sec)

mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/lng_synonym/lng_synonym.txt' INTO TABLE lng_synonym FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
```



```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/lng_synonym/lng_synonym.txt' INTO TABLE lng_synonym FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 1004 rows affected, 6 warnings (0.00 sec)
Records: 1004 Deleted: 0 Skipped: 0 Warnings: 6
```

```
mysql> CREATE TABLE lng_thl (lng_thl VARCHAR(1000) , PRIMARY KEY (lng_thl));
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/lng_thl/lng_thl.txt' INTO TABLE lng_thl FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 1064 rows affected (0.01 sec)
Records: 1064 Deleted: 0 Skipped: 0 Warnings: 0
```

```
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/lng_st/lng_st.txt' INTO TABLE lng_st FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 9 rows affected (0.00 sec)
Records: 9 Deleted: 0 Skipped: 0 Warnings: 0
```

```
mysql> CREATE TABLE lng_thp (lng_thp VARCHAR(100) , PRIMARY KEY (lng_thp));
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/lng_thp/lng_thp.txt' INTO TABLE lng_thp FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 782 rows affected (0.00 sec)
Records: 782 Deleted: 0 Skipped: 0 Warnings: 0
```

```
mysql> CREATE TABLE rel_code_name (lng_code VARCHAR(60) NOT NULL, lng_name VARCHAR(1000), IDKey INT NOT NULL AUTO_INCREMENT, PRIMARY KEY (IDKey));
Query OK, 0 rows affected (0.00 sec)
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/rel_code_name/rel_code_name.txt' INTO TABLE rel_code_name FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 1252 rows affected, 1252 warnings (0.00 sec)
Records: 1252 Deleted: 0 Skipped: 0 Warnings: 1252
```

```
mysql> CREATE TABLE rel_code_synonym (lng_synonym VARCHAR(1000), lng_code VARCHAR(60), IDKey INT NOT NULL AUTO_INCREMENT, PRIMARY KEY (IDKey));
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/rel_code_synonym/rel_code_synonym.txt' INTO TABLE rel_code_synonym FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 15281 rows affected, 15283 warnings (0.01 sec)
Records: 15281 Deleted: 0 Skipped: 0 Warnings: 2
```

```
mysql> CREATE TABLE rel_code_thl ( lng_code VARCHAR(60), lng_thl VARCHAR(1000), IDKey INT NOT NULL AUTO_INCREMENT, PRIMARY KEY (IDKey));
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/rel_code_thl/rel_code_thl.txt' INTO TABLE rel_code_thl FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 1249 rows affected, 1249 warnings (0.00 sec)
Records: 1249 Deleted: 0 Skipped: 0 Warnings: 0
```

```
mysql>
```

```
mysql> CREATE TABLE lng_id ( lng_code VARCHAR(60), a_lng_lat VARCHAR(1000), a_lng_lng VARCHAR(1000), lng_uri VARCHAR(1000), IDKey INT NOT NULL AUTO_INCREMENT, PRIMARY KEY (IDKey));
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/lng_id/lng_id.txt' INTO TABLE lng_id FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 1209 rows affected, 1209 warnings (0.00 sec)
Records: 1209 Deleted: 0 Skipped: 0 Warnings: 0
```

```
mysql>
```

```
mysql> CREATE TABLE rel_code_thp (lng_code VARCHAR(60), lng_thp VARCHAR(1000), IDKey INT NOT NULL AUTO_INCREMENT, PRIMARY KEY (IDKey));
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/rel_code_thp/rel_code_thp.txt' INTO TABLE rel_code_thp FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 1249 rows affected, 1249 warnings (0.00 sec)
Records: 1249 Deleted: 0 Skipped: 0 Warnings: 0
```

```
mysql> CREATE TABLE rel_code_st (lng_code VARCHAR(60), lng_st VARCHAR(1000), IDKey INT NOT NULL AUTO_INCREMENT, PRIMARY KEY (IDKey));
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> LOAD DATA INFILE '/home/training/workspace/Ass2/src/rel_code_st/rel_code_st.txt' INTO TABLE rel_code_st FIELDS TERMINATED BY '\t' ENCLOSED BY '"' LINES TERMINATED BY '\n';
Query OK, 1277 rows affected, 1277 warnings (0.00 sec)
Records: 1277 Deleted: 0 Skipped: 0 Warnings: 0
```

5.2. List all statements of this task for Hive:

First, we need to put the text all the txt files into HDFS.

```
hive> CREATE TABLE indiginous.lng_name (lng_name string) ROW FORMAT DELIMITED FI
ELDS TERMINATED BY '\t' LINES TERMINATED BY '\n' STORED AS TEXTFILE;
OK
Time taken: 0.106 seconds
hive> LOAD DATA INPATH 'lng_name.txt' INTO TABLE indiginous.lng_name;
Loading data to table indiginous.lng_name
OK
Time taken: 0.208 seconds

hive> CREATE TABLE indiginous.lng_synonym (lng_synonym string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATE
LINES TERMINATED BY '\n' STORED AS TEXTFILE;
OK
Time taken: 0.086 seconds

LOAD DATA INPATH 'lng_synonym.txt' INTO TABLE indiginous.lng_synonym;
Loading data to table indiginous.lng_synonym
OK
Time taken: 0.892 seconds

hive> CREATE TABLE indiginous.lng_thl (lng_thl string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n
' STORED AS TEXTFILE;
OK
Time taken: 0.032 seconds

LOAD DATA INPATH 'lng_thl.txt' INTO TABLE indiginous.lng_thl;
Loading data to table indiginous.lng_thl
OK
Time taken: 0.151 seconds

hive> CREATE TABLE indiginous.rel_code_name (lng_code string, lng_name string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
LINES TERMINATED BY '\n' STORED AS TEXTFILE;
OK
Time taken: 0.034 seconds

LOAD DATA INPATH 'rel_code_name.txt' INTO TABLE indiginous.rel_code_name;
Loading data to table indiginous.rel_code_name
OK
Time taken: 0.101 seconds

hive> CREATE TABLE indiginous.rel_code_synonym (lng_code string, lng_synonym string) ROW FORMAT DELIMITED FIELDS TERMINATED B
Y '\t' LINES TERMINATED BY '\n' STORED AS TEXTFILE;
hive> CREATE TABLE indiginous.rel_code_synonym (lng_code string, lng_synonym string) ROW FORMAT DELIMITED FIELDS TERMINATED B

LOAD DATA INPATH 'rel_code_synonym.txt' INTO TABLE indiginous.rel_code_synonym;
Loading data to table indiginous.rel_code_synonym
OK
Time taken: 0.117 seconds

hive> CREATE TABLE indiginous.rel_code_thl (lng_code string, lng_thl string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' L
INES TERMINATED BY '\n' STORED AS TEXTFILE;
OK
Time taken: 0.034 seconds
hive> CREATE TABLE indiginous.rel_code_thl (lng_code string, lng_thl string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' L

LOAD DATA INPATH 'rel_code_thl.txt' INTO TABLE indiginous.rel_code_thl;
Loading data to table indiginous.rel_code_thl
OK

LOAD DATA INPATH 'lng_id.txt' INTO TABLE indiginous.lng_id;
Loading data to table indiginous.lng_id
OK
Time taken: 0.105 seconds
hive>

hive> CREATE TABLE indiginous.rel_code_thp (lng_code string, lng_thp string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' L
INES TERMINATED BY '\n' STORED AS TEXTFILE;
OK
Time taken: 0.035 seconds
```

```

LOAD DATA INPATH 'rel_code_thp.txt' INTO TABLE indiginous.rel_code_thp;
Loading data to table indiginous.rel_code_thp
OK
Time taken: 0.116 seconds

hive> CREATE TABLE indiginous.rel_code_st (lng_code string, lng_st string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n' STORED AS TEXTFILE;
OK
Time taken: 0.058 seconds
hive> CREATE TABLE indiginous.rel_code_st (lng_code string, lng_st string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n' STORED AS TEXTFILE;
LOAD DATA INPATH 'rel_code_st.txt' INTO TABLE indiginous.rel_code_st;
Loading data to table indiginous.rel_code_st
OK
Time taken: 0.097 seconds

hive> CREATE TABLE indiginous.lng_thp (lng_thp string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n' STORED AS TEXTFILE;
OK
Time taken: 0.043 seconds
hive> CREATE TABLE indiginous.lng_thp (lng_thp string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n' STORED AS TEXTFILE;
LOAD DATA INPATH 'lng_thp.txt' INTO TABLE indiginous.lng_thp;
Loading data to table indiginous.lng_thp
OK

hive> CREATE TABLE indiginous.lng_st (lng_st string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n' STORED AS TEXTFILE;
OK
Time taken: 0.035 seconds
hive> CREATE TABLE indiginous.lng_st (lng_st string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n' STORED AS TEXTFILE;
LOAD DATA INPATH 'lng_st.txt' INTO TABLE indiginous.lng_st;
Loading data to table indiginous.lng_st
OK
Time taken: 0.105 seconds
hive>

```

6. Show the detailed information of table rel_code_st in MySQL.

6.1. Statement:

Ans: `DESCRIBE rel_code_st;`

6.2. Result screenshot:

```

Database changed
mysql> DESCRIBE rel_code_st;
+-----+-----+-----+-----+-----+-----+
| Field | Type          | Null | Key | Default | Extra          |
+-----+-----+-----+-----+-----+-----+
| lng_code | varchar(60)   | YES  |     | NULL    |                |
| lng_st   | varchar(1000) | YES  |     | NULL    |                |
| IDKey    | int(11)       | NO   | PRI | NULL    | auto_increment |
+-----+-----+-----+-----+-----+-----+
3 rows in set (0.00 sec)

```

7. Collect the results containing “lng_code”, “lng_name”, and “lng_st” of tuples whose lng_name starts with the upper case “D”. Show bot statements and the last page screenshot of your query result in MySQL and Hive.

7.1. MySQL query statement:

```

mysql> SELECT rel_code_name.lng_code, rel_code_name.lng_name, rel_code_st.lng_st FROM rel_code_name INNER JOIN rel_code_st ON
rel_code_name.lng_code = rel_code_st.lng_code WHERE lng_name LIKE 'D%';

```

7.2. MySQL result screenshot:

Ans:

```

| N209 | Dbi          | NT |
| N30  | Djowei       |    |
| N38  | Dagoman      | NT |
| N60  | Dalabon      | NT |
| N84  | Djarawala    |    |
| N94.1 | Djinang      | NT |
| N96  | Dhaygurrurr  | NT |
| N97  | Djinba^      | NT |
| S20  | Dhauwurd Wurrug^ | VIC |
| S22  | Djadjala     | VIC |
| S26  | Djab Wurrug^ | VIC |
| S28  | Dadi Dadi    | NSW |
| S28  | Dadi Dadi    | VIC |
| S31.1 | Dja Dja Wurrug | VIC |
| S44  | Dhudhuroa    | VIC |
| S51  | Dyirringaã   | NSW |
| S53  | Dhurga       | NSW |
| S56  | Dharamba     | NSW |
| S59  | Dharawal     | NSW |
| S64  | Dharug       | NSW |
| S65  | Darkinyung   | NSW |
| S73  | Djargurd Wurrong | VIC |
| S95  | Djilamatang  |    |
| W16  | Damala       |    |
| W48  | Djungundja   |    |
| W65  | Djiraly      |    |
| Y106 | Djabugay     | QLD |
| Y116 | Djandjandji  |    |
| Y123 | DYIRBAL      | QLD |
| Y124 | Djiru        | QLD |
| Y167 | Dhalundhirr  | QLD |
| Y221 | Di:ru        |    |
| Y227 | Daru         |    |
| Y235 | Dulgubarra Mamu |    |
| Y6   | Djagaraga    | QLD |
+-----+-----+-----+
76 rows in set (0.00 sec)

```

7.3. Hive query statement:

```

hive> SELECT rel_code_name.lng_code, rel_code_name.lng_name, rel_code_st.lng_st FROM rel_code_name INNER JOIN rel_code_st ON
rel_code_name.lng_code = rel_code_st.lng_code WHERE lng_name LIKE 'D%';

```

7.4. Hive result screenshot:

Ans:

```

N202  Djangu  NT
N206  Djadiwitjibi  NT
N209  Dbi      NT
N30   Djowei
N38   Dagoman      NT
N60   Dalabon      NT
N84   Djarawala
N94.1 Djinang      NT
N96   Dhaygurrurr  NT
N97   Djinba^ NT
S20   Dhauwurd Wurrug^  VIC
S22   Djadjala     VIC
S26   Djab Wurrug^  VIC
S28   Dadi Dadi    NSW
S28   Dadi Dadi    VIC
S31.1 Dja Dja Wurrug  VIC
S44   Dhudhuroa    VIC
S51   Dyirringaã   NSW
S53   Dhurga       NSW
S56   Dharamba     NSW
S59   Dharawal     NSW
S64   Dharug       NSW
S65   Darkinyung   NSW
S73   Djargurd Wurrong  VIC
S95   Djilamatang
W16   Damala
W48   Djungundja
W65   Djiraly
Y106  Djabugay     QLD
Y116  Djandjandji
Y123  DYIRBAL      QLD
Y124  Djiru       QLD
Y167  Dhalundhirr  QLD
Y221  Di:ru
Y227  Daru
Y235  Dulgubarra Mamu
Y6    Djagaraga    QLD
Time taken: 15.312 seconds
hive>

```

8. Collect the results containing "lng_code", "lng_name", "lng_st", "a_lng_lat", and "a_lng_lng" of tuples whose lng_synonym contains "Kerama". Show both statements and the last page screenshot of your query result in MySQL and Hive.

8.1. How many tuples are retrieved at the end?

Ans: 9 tuples are retrieved at the end.

8.2. MySQL query statement:

```
SELECT lng_id.lng_code, rel_code_name.lng_name, rel_code_st.lng_st, lng_id.a_lng_lat, lng_id.a_lng_lng FROM lng_id INNER JOIN rel_code_name ON lng_id.lng_code = rel_code_name.lng_code INNER JOIN rel_code_st ON lng_id.lng_code = rel_code_st.lng_code INNER JOIN rel_code_synonym ON lng_id.lng_code = rel_code_synonym.lng_code WHERE lng_synonym LIKE '%Kerama%';
```

8.3. MySQL result screenshot:

Ans:

lng_code	lng_name	lng_st	a_lng_lat	a_lng_lng
W31	Yarnarri			
W36	Kurrama	WA	-22.3454093	117.1252561
W49	Jadira	WA	-21.84552987	116.1884783
Y121	Ngadjon	QLD	-17.36377694	145.7047546
Y122	Mamu	QLD	-17.63998402	145.7485185
Y123	DYIRBAL	QLD	-17.74022656	145.661334
Y123	JIRRBAL	QLD	-17.74022656	145.661334
Y124	Djiru	QLD	-17.90722366	146.0521375
Y126	Gulngay	QLD	-17.96177304	145.8525333

9 rows in set (0.10 sec)

8.4. Hive query statement:

```
hive> SELECT lng_id.lng_code, rel_code_name.lng_name, rel_code_st.lng_st, lng_id.a_lng_lat, lng_id.a_lng_lng FROM lng_id INNER JOIN rel_code_name ON lng_id.lng_code = rel_code_name.lng_code INNER JOIN rel_code_st ON lng_id.lng_code = rel_code_st.lng_code INNER JOIN rel_code_synonym ON lng_id.lng_code = rel_code_synonym.lng_code WHERE lng_synonym LIKE '%Kerama%';
```

8.5. Hive result screenshot:

Ans:

```
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapred.reduce.tasks=<number>
Starting Job = job_202310271501_0011, Tracking URL = http://0.0.0.0:50030/jobdetails.jsp?jobid=job_202310271501_0011
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_202310271501_0011
Hadoop job information for Stage-1: number of mappers: 4; number of reducers: 1
023-10-27 16:37:01,491 Stage-1 map = 0%, reduce = 0%
023-10-27 16:37:06,527 Stage-1 map = 25%, reduce = 0%, Cumulative CPU 0.55 sec
023-10-27 16:37:07,532 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 1.45 sec
023-10-27 16:37:08,536 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 1.45 sec
023-10-27 16:37:09,542 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 1.45 sec
023-10-27 16:37:10,559 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 1.45 sec
023-10-27 16:37:11,572 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 1.45 sec
023-10-27 16:37:12,576 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 1.45 sec
023-10-27 16:37:13,582 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.57 sec
023-10-27 16:37:14,589 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.71 sec
023-10-27 16:37:15,596 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.71 sec
023-10-27 16:37:16,600 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.71 sec
mapReduce Total cumulative CPU time: 3 seconds 710 msec
Ended Job = job_202310271501_0011
mapReduce Jobs Launched:
Job 0: Map: 4 Reduce: 1 Cumulative CPU: 3.71 sec HDFS Read: 368521 HDFS Write: 346 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 710 msec
OK
31 Yarnarri
36 Kurrama WA -22.3454093 117.1252561
49 Jadira WA -21.84552987 116.1884783
121 Ngadjon QLD -17.36377694 145.7047546
122 Mamu QLD -17.63998402 145.7485185
123 DYIRBAL QLD -17.74022656 145.661334
123 JIRRBAL QLD -17.74022656 145.661334
124 Djiru QLD -17.90722366 146.0521375
126 Gulngay QLD -17.96177304 145.8525333
Time taken: 17.716 seconds
hive>
```